

# Homework 1:

---

## 1. Problem Overview

The objective is to develop a multimodal AI agent capable of processing multiple supermarket receipt images to answer specific financial queries:

1. **Total Expenditure:** Summing the final amounts charged across all receipts.
  2. **Pre-discount Total:** Calculating what the total would have been without any applied savings.
  3. **Query Routing:** Identifying and rejecting out-of-domain or irrelevant questions.
- 

## 2. Agentic Architecture & Solution

The solution is built using a **Router-based Agentic Workflow** leveraging LangChain and Gemini 2.5 Flash. Instead of a single monolithic prompt, the system uses a modular pipeline:

### 2.1 Multimodal Input Handling

Images are pre-processed into **Base64 encoded Data URLs**. This allows the model to receive a list of image objects alongside text instructions, enabling the Gemini model to perform cross-image reasoning.

### 2.2 The Router Pattern

To satisfy the requirement of rejecting irrelevant queries, a **Router Chain** acts as the "brain" of the agent. It classifies the intent into three categories: `cost`, `original`, or `irrelevant`.

- **Logic:** The LLM evaluates the user query first.
- **Safety:** If categorized as `irrelevant`, the agent triggers a specific branch that politely declines the request, ensuring the system remains within its functional guardrails.

## 2.3 Structured Data Extraction

Rather than asking the LLM to do "mental math" directly on images (which often leads to hallucinations), I implemented **Structured Output Extraction**.

- **Pydantic Schema:** I defined a `ReceiptData` model and a `Discount` model.
  - **Extraction:** The agent extracts raw financial data into these structured JSON objects.
  - **Computational Accuracy:** Once extracted, the math is performed using Python functions (`sum()` and `calculate_original_price`) rather than LLM generation, ensuring 100% arithmetic precision.
- 

## 3. Implementation Details

### 3.1 Data Models

```
class Discount(BaseModel):  
    item_name: str  
    discount_reason: str  
    amount: float  
  
class ReceiptData(BaseModel):  
    subtotal: float  
    discounts: List[Discount]
```

### 3.2 The Decision Branch

The workflow uses a `RunnableBranch` to direct the flow based on the router's output:

- **Branch A (Total Spend):** Extracts `subtotal` from all images and sums them.
  - **Branch B (Original Price):** Extracts `subtotal` and all `discounts`, then adds them back together for each receipt before summing the grand total.
  - **Branch C (Fallback):** Returns a standard "irrelevant query" message.
-

## 4. Evaluation and Results

The agent was tested against seven receipt images. The evaluation criteria required the calculated totals to be within **\$2.00** of the ground truth.

Query Type	Ground Truth Total	Agent Result	Status
Total Spent (Query 1)	\$1974.30	\$1974.64	Passed
Original Price (Query 2)	\$2348.20	\$2348.39	Passed
Out-of-Domain	N/A	"I'm sorry..."	Correctly Rejected

## 5. Conclusion

By using an **Agentic Router** combined with **Structured Extraction**, the solution overcomes the typical limitations of LLMs in arithmetic. By separating the *interpretation* of the image from the *calculation* of the values, the agent provides a robust and reliable financial tool for business applications.